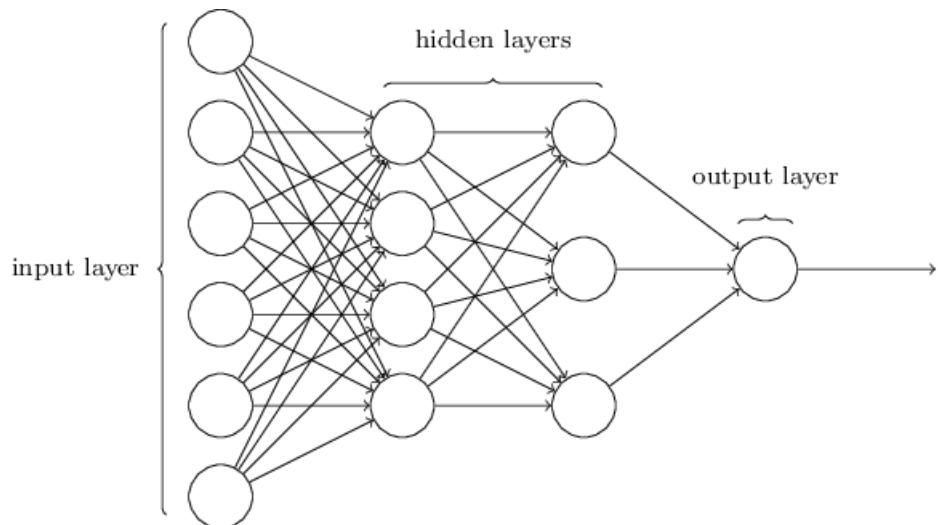


```
In [1]: from IPython.display import Image
```

## Neural Networks and Deep Learning

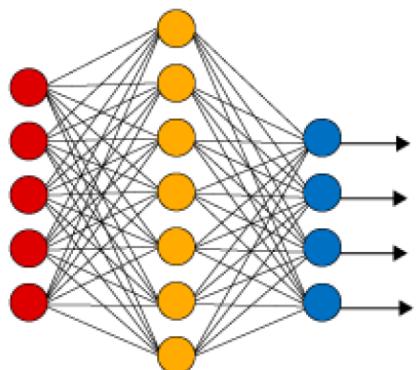
What is a Neural Network?



Input Layer and middle layer is densely conneted.

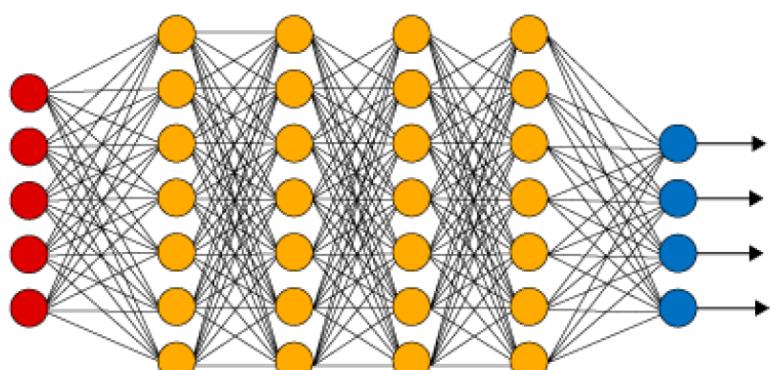
How a neural network looks?

Simple Neural Network



● Input Layer

Deep Learning Neural Network

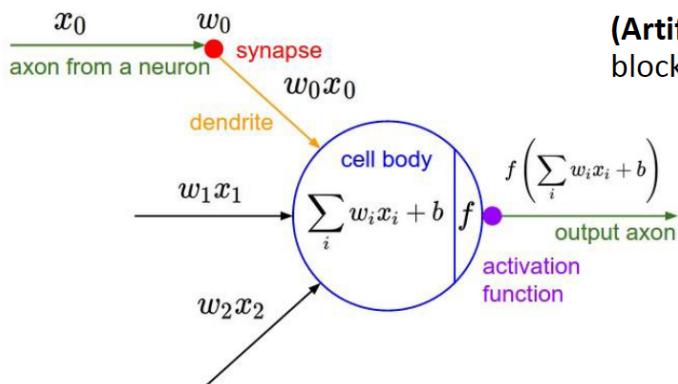


● Hidden Layer

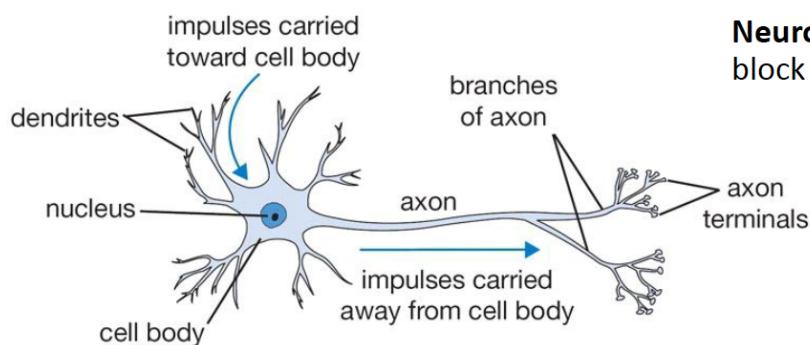
● Output Layer

**Universality:** For any arbitrary function  $f(x)$ , there exists a neural network that closely approximate it for any input  $x$

## Neuron: Biological Inspiration for Computation



**(Artificial) Neuron:** computational building block for the “neural network”



**Neuron:** computational building block for the brain

### Supervised Learning with Neural Network

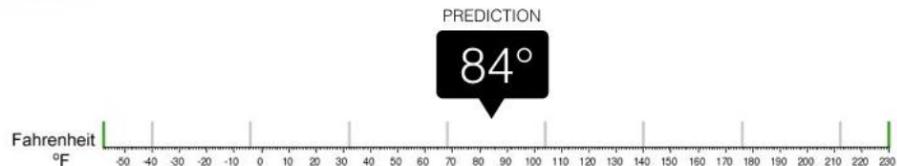
Input (X)	Output (y)	Application	
Home Features	Price	Real Estate	Standard Neural Networks
Ad, User info	Click on Ad? (0/1)	Online Advertising	Standard Neural Networks
Image	Object (1,2,---,100)	Photo Tagging	CNN
Audio	Text Transcript	Speech Recognition	RNN
English	Chinese	Machine Translation	RNN
Image, Radar Info	Position of other cars	Autonomous Driving	Custom/Hybrid

### Regression vs Classification



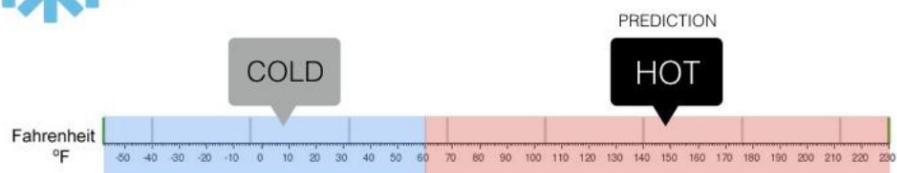
# Regression

What is the temperature going to be tomorrow?

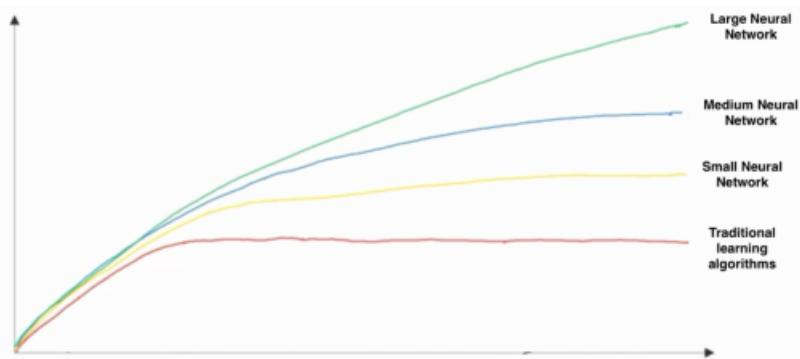


## Classification

## Will it be Cold or Hot tomorrow?



## Why is DeepLearning taking off?



The vertical axes of the diagram you can see the performance of an algorithm (e.g. it's prediction accuracy) and at the horizontal axes you can see the amount of data (Labelled Data).

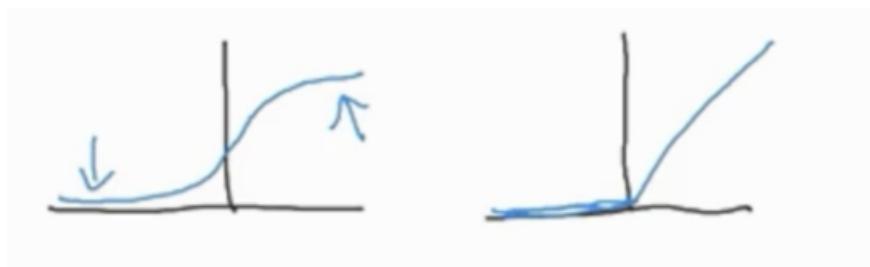
You can also see that the performance of traditional learning algorithms (logistic Regression, SVM's etc.) increases at the beginning with an increase of the amount of data but that it plateaus at a certain level and stops improving its performance.

The thing is that we have accumulated huge amounts of data over the last decades where our traditional learning algorithms can't take advantage of, which is where Deep Learning comes into play.

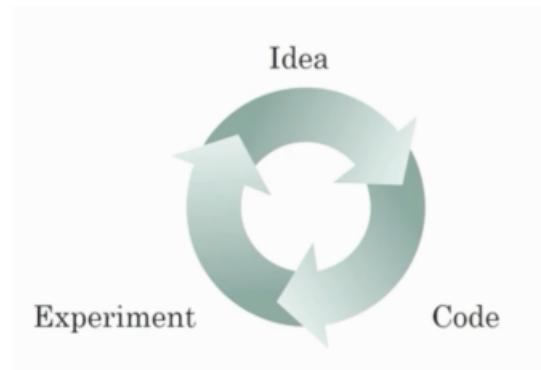
Large Neural Networks (e.g. Deep Learning) are getting better and better the more data you put into them. Andrew NG, a leading AI scientist, said that the three main forces which improve Neural Networks are:

1. Data
2. Computation
3. Algorithms

The recent breakthroughs in the development of algorithms are mostly due to making them run much faster than before, which makes it possible to use more and more data. For an example, a big advancement came from switching from a Sigmoid function (left picture) to a rectified-Linear-Unit function (right picture).



The other reason why fast computation is important is that, the below cycle must be faster.



**Bringing more data to a model is almost always beneficial.**

**Deep Learning approaches improve with more data**

## Deep Learning Representation

## 2 Deep Learning representations

For representations:

- nodes represent inputs, activations or outputs
- edges represent weights or biases

Here are several examples of Standard deep learning representations

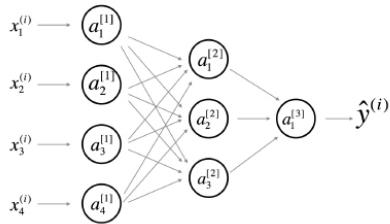


Figure 1: Comprehensive Network: representation commonly used for Neural Networks. For better aesthetic, we omitted the details on the parameters ( $w_{ij}^{[l]}$  and  $b_i^{[l]}$  etc...) that should appear on the edges

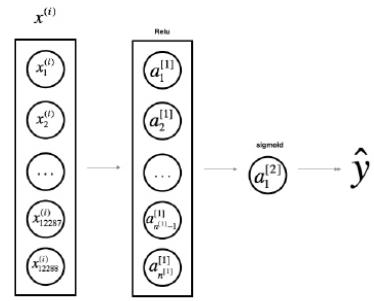


Figure 2: Simplified Network: a simpler representation of a two layer neural network, both are equivalent.

## Deep Learning

### • What is it:

Extract useful patterns from data.

### • How:

Neural network + optimization

### • How (Practical):

Python + TensorFlow & friends

### • Hard Part:

Good Questions + Good Data

### • Why now:

Data, hardware, community, tools, investment

### • Where do we stand?

Most big questions of intelligence have not been answered nor properly formulated

### Exciting progress:

- Face recognition
- Image classification
- Speech recognition
- Text-to-speech generation
- Handwriting transcription
- Machine translation
- Medical diagnosis
- Cars: drivable area, lane keeping
- Digital assistants
- Ads, search, social recommendations
- Game playing with deep RL

## Deep Learning Tools

1. CUDA
2. Theno
3. Caffe
4. Tensorflow 0.1 (2015)
5. Pytorch 0.1 (2017)
6. **Tensorflow 1.0 (2017)**
7. **Pytorch 1.0 (2017)**
8. **Tensorflow 2.0 (2019)**

## What is Tensorflow?

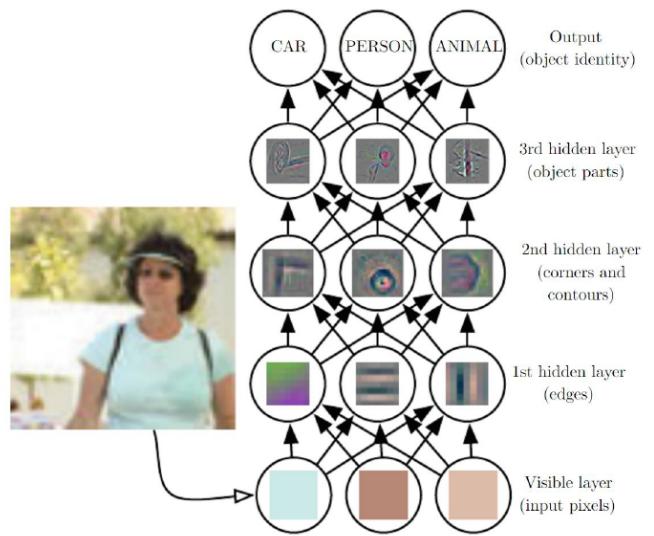
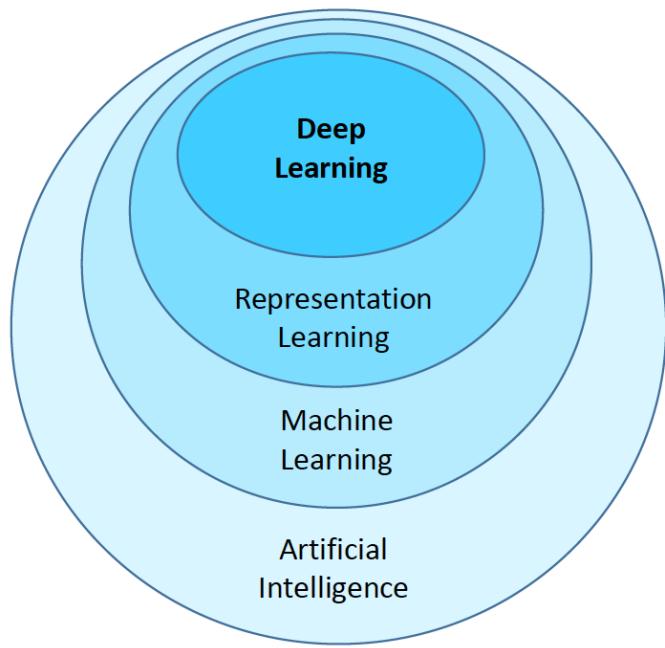
- **What is it:** Deep Learning Library (*and more*)
  - **Facts:** Open Source, Python, Google
- **Community:**
  - 117,000+ GitHub stars
  - TensorFlow.org: Blogs, Documentation, DevSummit, YouTube talks
- **Ecosystem:**
  - **Keras:** high-level API
  - **TensorFlow.js:** in the browser
  - **TensorFlow Lite:** on the phone
  - **Colaboratory:** in the cloud
  - **TPU:** optimized hardware
  - **TensorBoard:** visualization
  - **TensorFlow Hub:** graph modules
- **Alternatives:** PyTorch, MXNet, CNTK

## Extras:

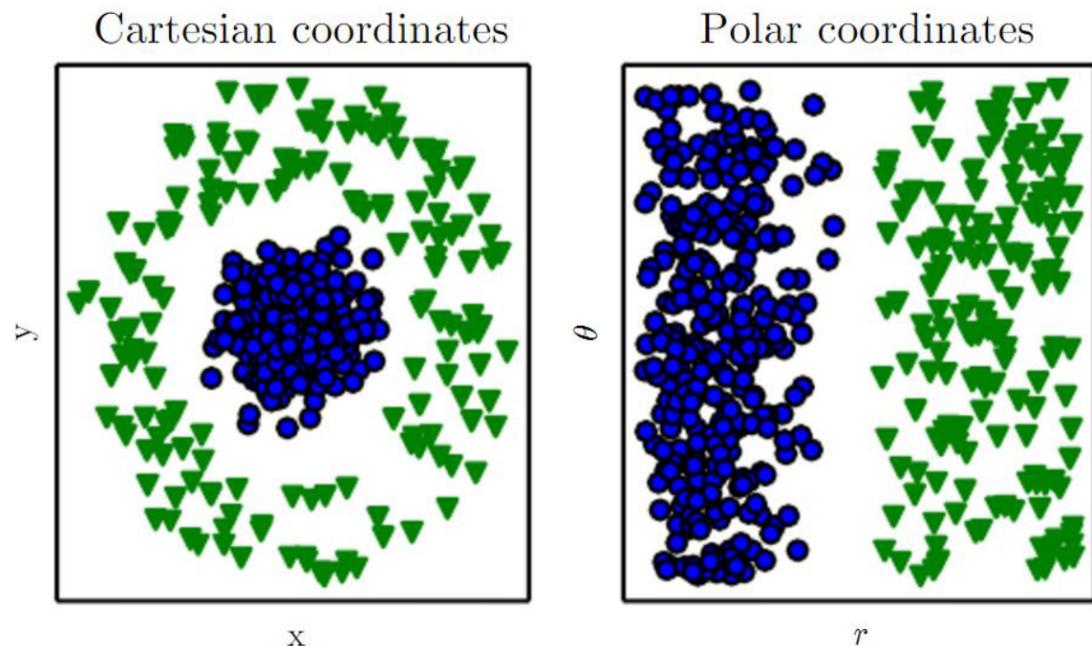
- Swift for TensorFlow
- TensorFlow Serving
- TensorFlow Extended (TFX)
- TensorFlow Probability
- Tensor2Tensor

## Deep Learning is Representation Learning

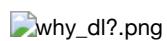
(aka Feature Learning)



# Representation Matters

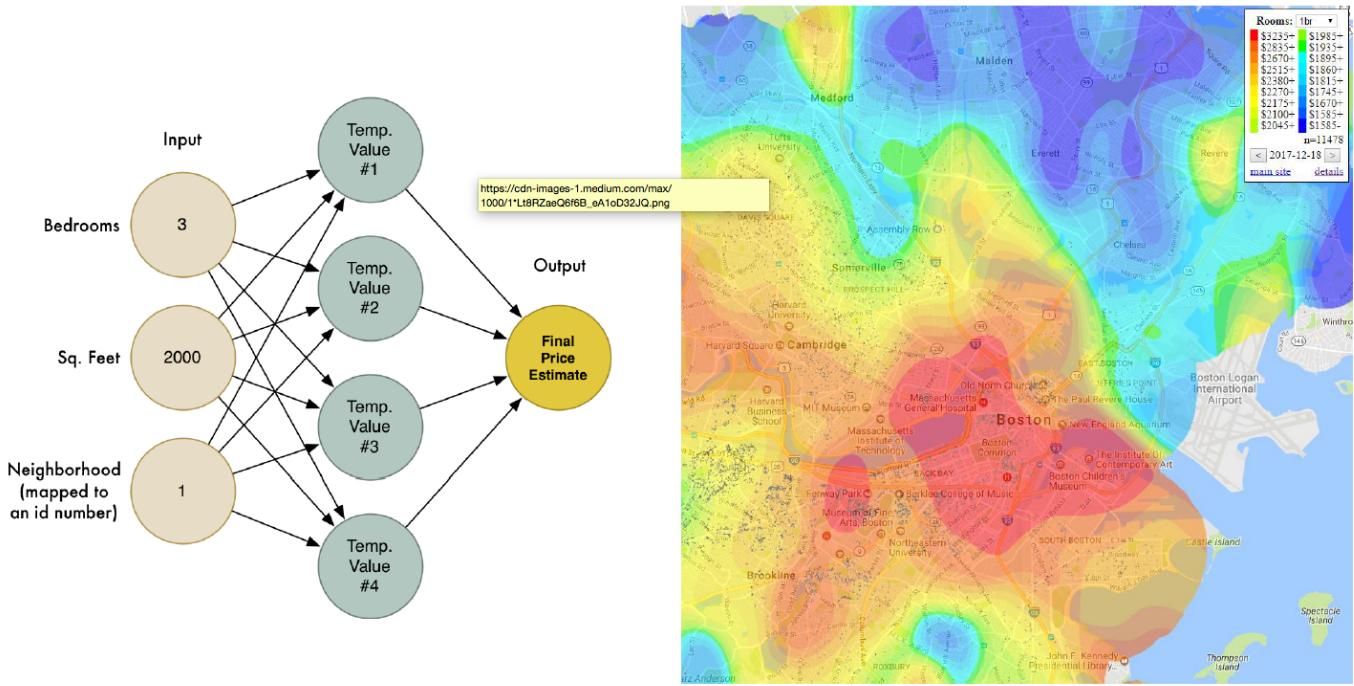


## WHY DEEP LEARNING?

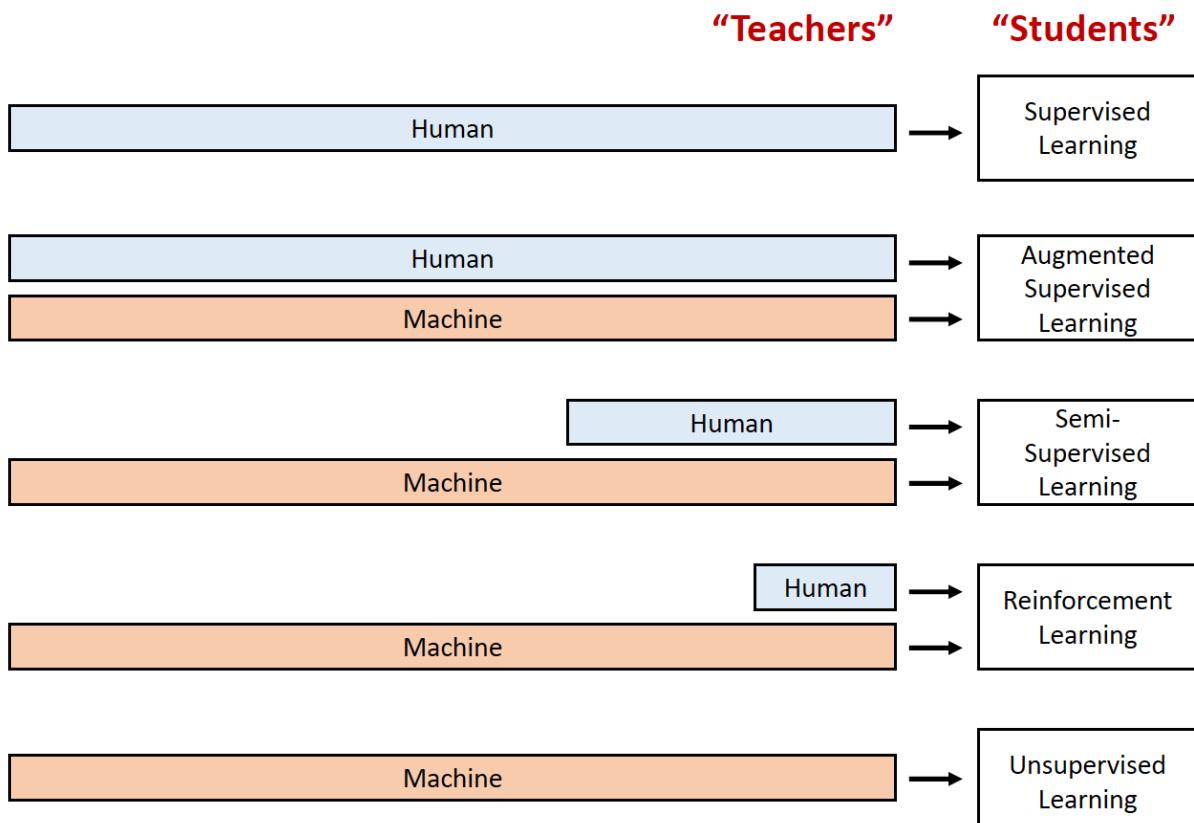


## What is a Neural Network?

By the Example of House Price Prediction



## DeepLearning for human and machine



## Data Augmentation

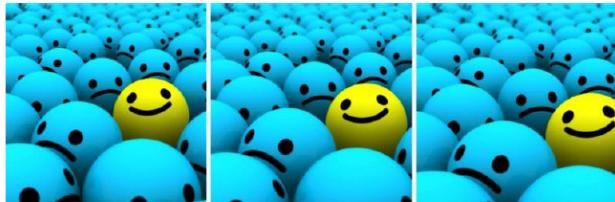
Crop:



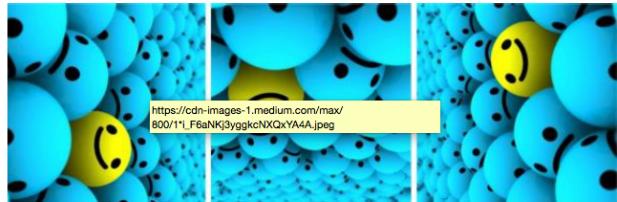
Flip:



Scale:



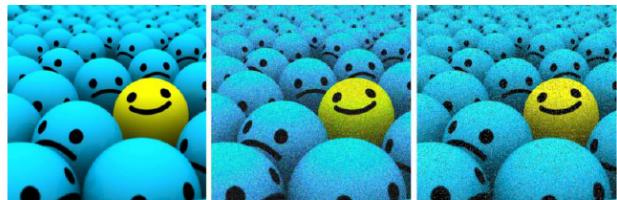
Rotate:



Translation:

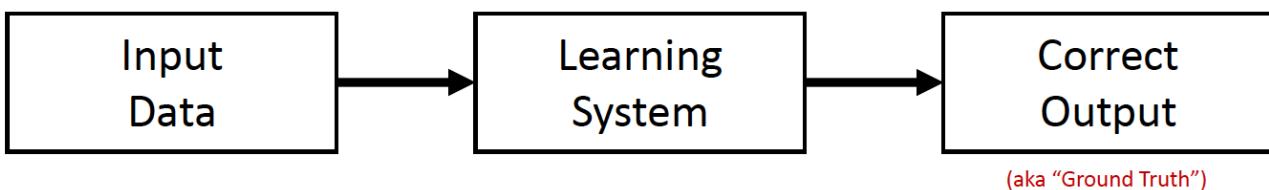


Noise:

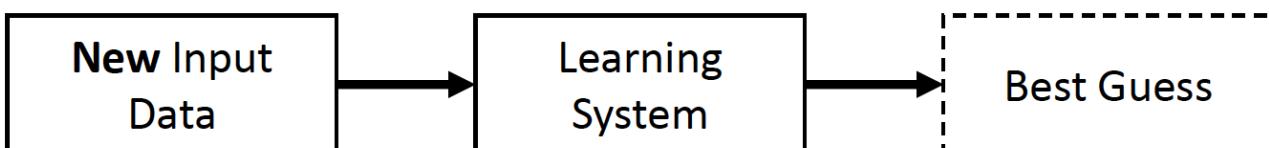


## Deep Learning: Training and Testing

### Training Stage:

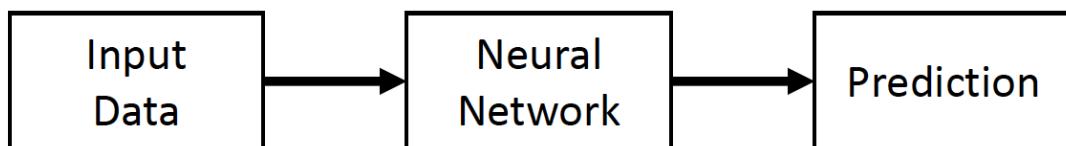


### Testing Stage:

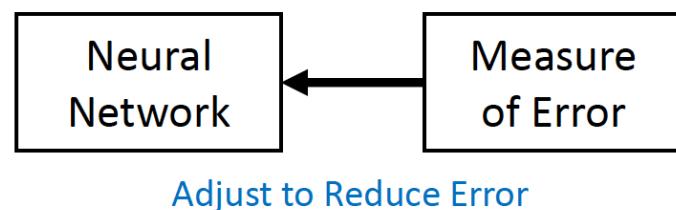


## How Neural Networks Learns : BackPropagation

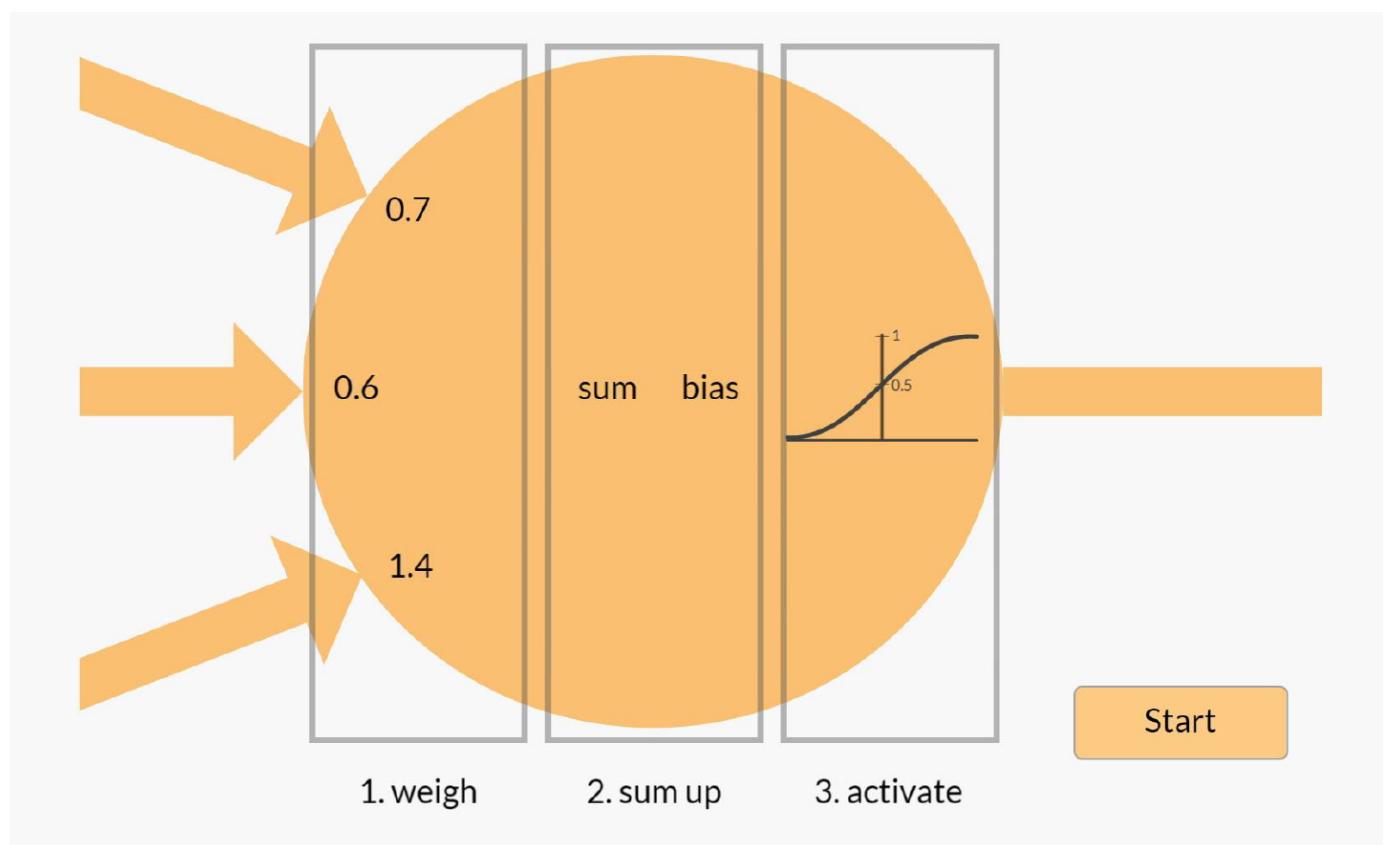
Forward Pass:



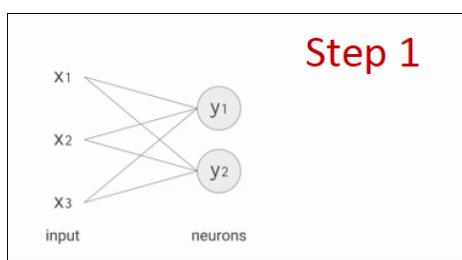
Backward Pass (aka Backpropagation):



Neuron : Forward pass



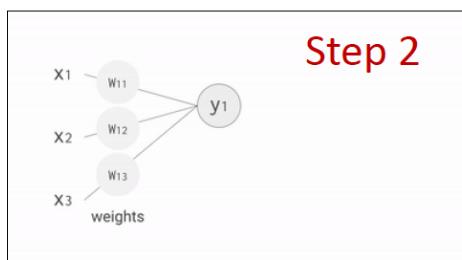
## Steps Involved in forward propagation!!



**Step 4**

$$y_1 = f(w_{11}x_1 + w_{12}x_2 + w_{13}x_3)$$

↑  
activation function



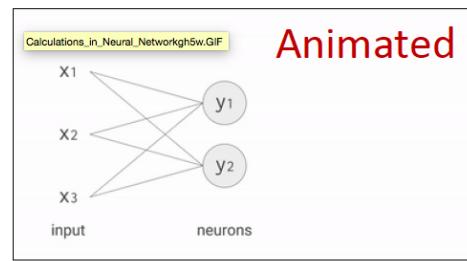
**Step 5**

$$y_1 = f(w_{11}x_1 + w_{12}x_2 + w_{13}x_3)$$

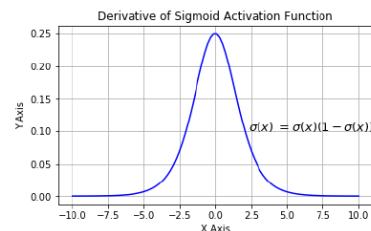
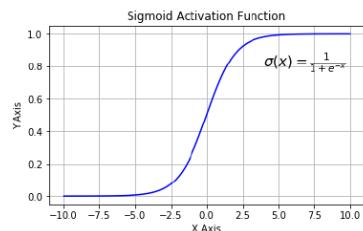
$$y_2 = f(w_{21}x_1 + w_{22}x_2 + w_{23}x_3)$$

**Step 3**

$$y_1 = w_{11}x_1 + w_{12}x_2 + w_{13}x_3$$

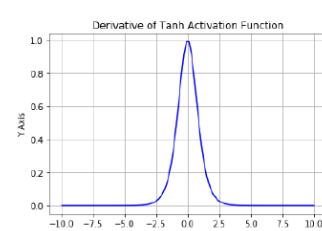
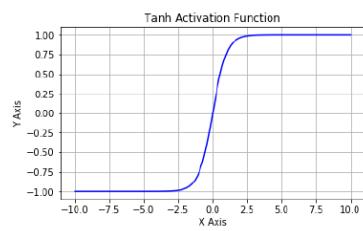


## Different types of Activation function



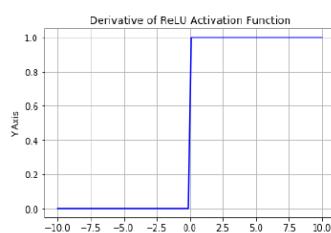
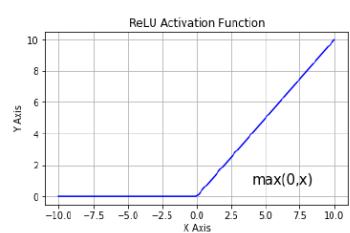
### Sigmoid

- Vanishing gradients
- Not zero centered



### Tanh

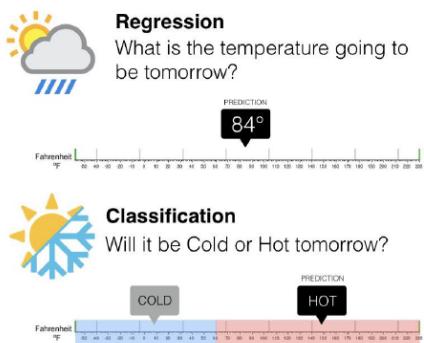
- Vanishing gradients



### ReLU

- Not zero centered

## Loss Function



- Loss function quantifies gap between prediction and ground truth
- For regression:
  - Mean Squared Error (MSE)
- For classification:
  - Cross Entropy Loss

### Mean Squared Error

$$MSE = \frac{1}{N} \sum (t_i - s_i)^2$$

Prediction  
↓  
Ground Truth

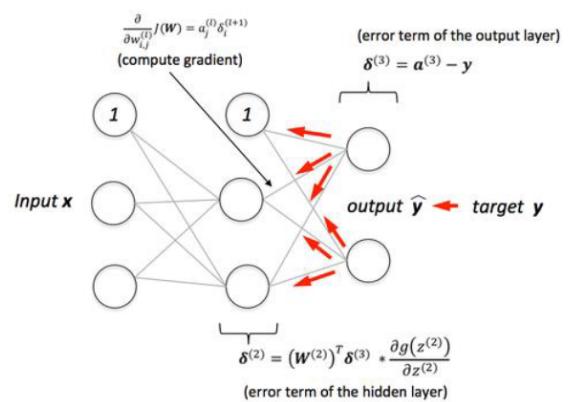
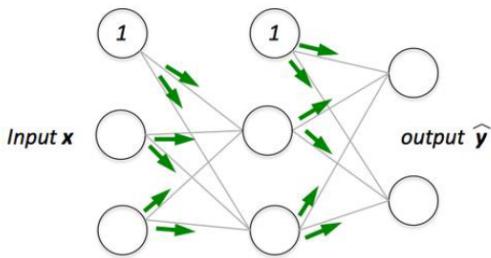
### Cross Entropy Loss

$$CE = - \sum_i t_i \log(s_i)$$

Classes  
↓  
Prediction  
↓  
Ground Truth {0,1}

## Steps Involved in

### Backpropagation



**Task:** Update the **weights** and **biases** to decrease **loss function**

Subtasks:

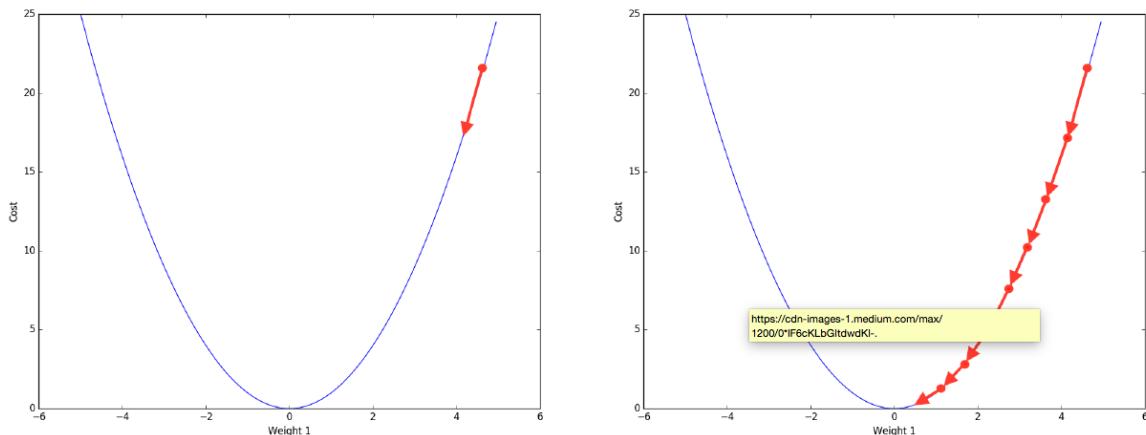
1. Forward pass to compute network output and “error”
2. Backward pass to compute gradients
3. A fraction of the weight’s gradient is subtracted from the weight.

↑  
Learning Rate

Numerical Method: **Automatic Differentiation**

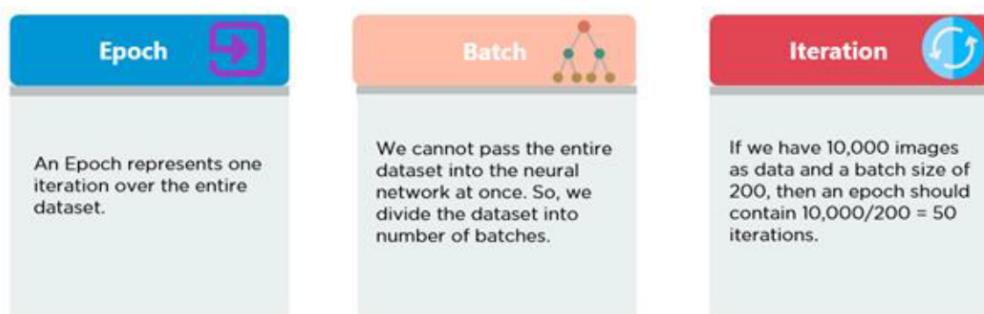
## Learning is an Optimizing Problem

Task: Update the **weights** and **biases** to decrease loss function



SGD: Stochastic Gradient Descent

## Mini-Batch Size

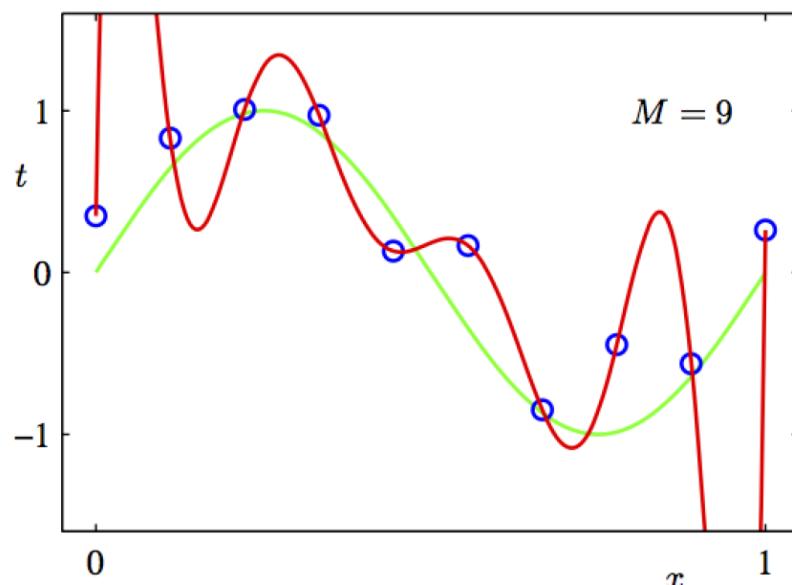


**Mini-Batch size:** Number of training instances the network evaluates per weight update step.

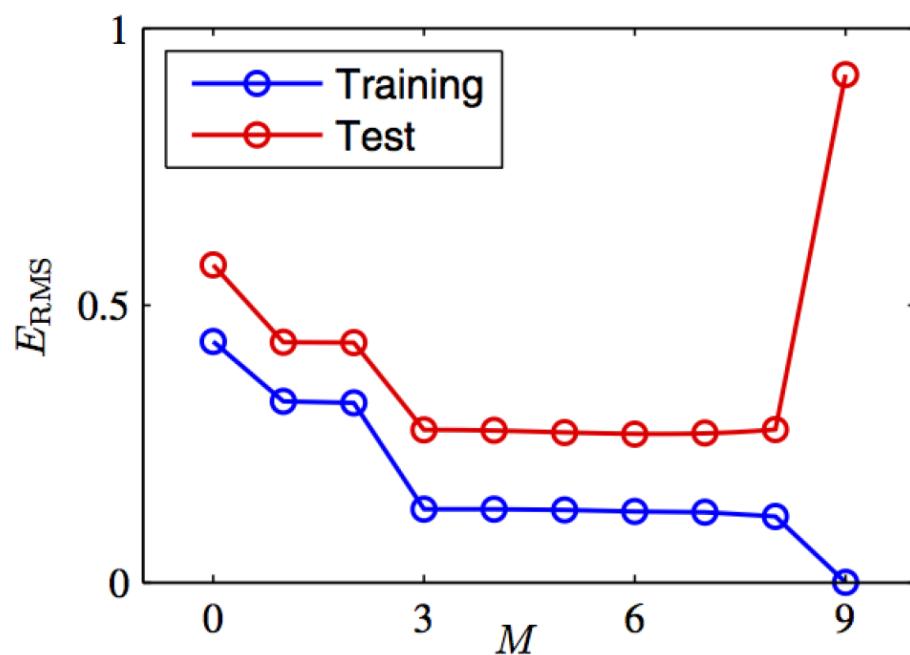
- Larger batch size = more computational speed
- Smaller batch size = (empirically) better generalization

## Overfitting

- Help the network **generalize** to data it hasn't seen.
- Big problem for **small datasets**.
- Overfitting example (a sine curve vs 9-degree polynomial):

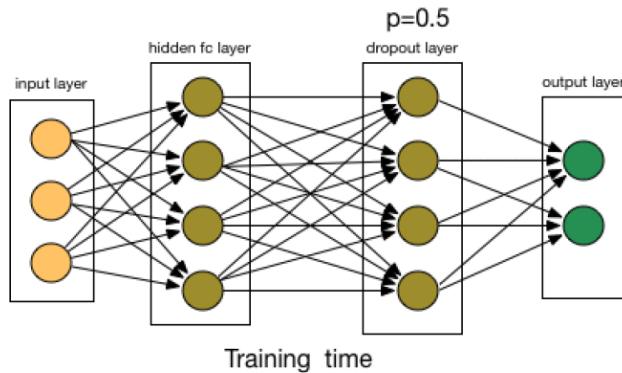


- Overfitting: The error decreases in the training set but increases in the test set.



---

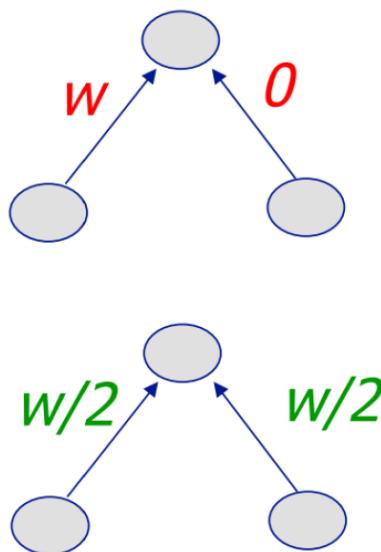
## Regularization : 1. Dropout



- **Dropout:** Randomly remove some nodes in the network (along with incoming and outgoing edges)
- Notes:
  - Usually  $p \geq 0.5$  ( $p$  is probability of keeping node)
  - Input layers  $p$  should be much higher (and use noise instead of dropout)
  - Most deep learning frameworks come with a dropout layer

## 2. Weight Penalty

# Regularization: Weight Penalty (*aka Weight Decay*)



- **L2 Penalty:** Penalize squared weights. Result:
  - Keeps weight small unless error derivative is very large.
  - Prevent from fitting sampling error.
  - Smoother model (output changes slower as the input change).
  - If network has two similar inputs, it prefers to put half the weight on each rather than all the weight on one.
- **L1 Penalty:** Penalize absolute weights. Result:
  - Allow for a few weights to remain large.

---

## Normalization

- Network Input Normalization
  - Example: Pixel to [0, 1] or [-1, 1] or according to mean and std.
- Batch Normalization (BatchNorm, BN)
  - Normalize hidden layer inputs to mini-batch mean & variance
  - Reduces impact of earlier layers on later layers
- Batch Renormalization (BatchRenorm, BR)
  - Fixes difference b/w training and inference by keeping a moving average asymptotically approaching a global normalization.
- Other options:
  - Layer normalization (LN) – conceived for RNNs
  - Instance normalization (IN) – conceived for Style Transfer
  - Group normalization (GN) – conceived for CNNs

---

## Neural Network Playground

<http://playground.tensorflow.org/> (<http://playground.tensorflow.org/>)

In [2]:

```
%%html
<iframe src="http://playground.tensorflow.org/" width="1000" height="600"></iframe>
```

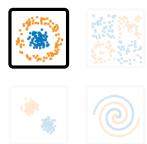
# Tinker With a Neural Network in Your Browser

## Don't Worry, You Can't Break It. We Promise

Epoch: 000,000      Learning rate: 0.03      Activation: Tanh      Regularization: None      R: 0

### DATA

Which dataset do you want to use?

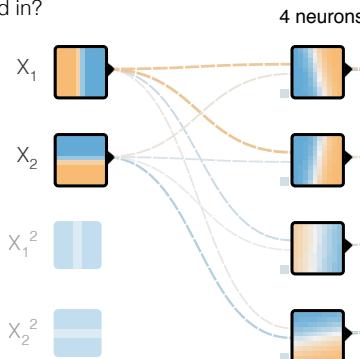


Ratio of training to test data: 50%

Noise: 0

### FEATURES

Which properties do you want to feed in?



### 2 HIDDEN LAYERS

Rfc\_msmisq\_ipc  
kgcb\_ugf\_t\_pnd\_e  
weights\*qfmul  
`wrf\_c\_ifgai\_lcqg  
mdif\_c\_jg cq,

### OUTPUT

Test loss 0.  
Training los:



In [ ]: